Technical diligence note: Composio (tool integrations for AI agents)

Practice note from public sources · Pavitra Seth · 8 Nov 2026

Why this company. Composio fits an India-focused AI fund's thesis: Indian founders building globally, in agent infrastructure. It is past seed, so this is a case study in *how* I would run technical diligence, not a recommendation. Registry figures were pulled by `src/package_health.py` on 8 Nov 2026; everything else is cited at the end.

1. Snapshot

- **What it does:** lets AI agents take actions in 1,000+ apps (Slack, GitHub, Salesforce, HubSpot, Gmail and more), with managed authentication (OAuth, API keys), tool search, context management and a sandboxed workbench.
- **Team:** founded 2022 by Soham Ganatra (CEO, IIT Bombay) and Karan Vaidya; HQ San Francisco. Inc42 lists the founding year as 2023, so check this.
- **Funding:** USD 25M Series A led by Lightspeed (USD 29M total). Existing investors: Elevation Capital and Together Fund. Angels include Guillermo Rauch (Vercel) and Dharmesh Shah (HubSpot).
- **Traction (company-reported):** 100,000+ developers and 200+ customers, including Glean.

2. What the public code and registries show

Architecture: the open-source SDK is a thin client. The Python `composio` package depends on a generated API client (`composio-client`), pydantic and JSON-schema tools. The npm package `@composio/core` depends on `@composio/client`. The value (auth, toolkits, execution) sits in the **hosted platform**, while the MIT-licensed SDK is distribution. So the moat is toolkit breadth, managed auth and reliability, not code anyone can copy. The org also publishes Helm charts, consistent with self-hosting for enterprise.

Shipping pace: fast, and recently re-platformed.

Package	Releases (total)	Last 90 days	Median days between releases (12m)
PyPI `composio` (new SDK, first release Dec 2024)	86	16	6
npm `@composio/core`	156	16	4
PyPI `composio-core` (old SDK, since Apr 2024)	223	~0	n/a (largely quiet since mid-2025)

The move from `composio-core` to `composio` looks like a v3 rewrite. **Diligence question:** what did the migration cost existing customers, and how are breaking changes handled?

Dependencies worth asking about. The core SDK lists the **OpenAI client** as a required dependency in both Python and JavaScript, even though the docs cover Gemini, LangChain/LangGraph, the Vercel

AI SDK and others. Real-time triggers run over **Pusher** (`pysher`, `pusher-js`), a third-party service: what happens to triggers if it degrades, and what does it cost at scale?

Attention is not usage. The product repo has ~30K GitHub stars, but the same org's curated list *awesome-claude-skills* has ~73K, more than the product. Composio is clearly good at distribution, but stars can't show usage. Usage needs active API keys, tool calls per account, paid conversion and net revenue retention.

3. Market and risks (my analysis)

- **Crowded layer:** agent integrations are also offered by Merge (Agent Handler), Zapier MCP, Pipedream MCP and Toolhouse.
- **Commoditisation:** MCP standardises how agents call tools, and model labs ship native connectors. Composio's defence has to be what's hard to copy: auth at enterprise scale, 1,000+ maintained toolkits, reliability, and the "learning" layer it is raising for.
- **Pricing:** four tiers metered on tool calls, premium tool calls and connected accounts. This comes from a competitor's blog, so I haven't verified the prices.

4. Questions I would ask the founders

1. **Reliability:** tool-call success rate and p95 latency per toolkit; how they detect an upstream API change before customers do.
2. **Maintenance cost:** share of toolkits generated automatically vs maintained by hand; engineers per 100 toolkits.
3. **Security:** how OAuth tokens are stored and scoped; audit logs; which certifications are complete.
4. **Unit economics:** gross margin per tool call; revenue concentration in the top 10 customers; net revenue retention by cohort.
5. **Dependencies:** why OpenAI is a hard dependency; plan if Pusher or a model provider changes terms.
6. **The learning layer:** what metric proves agents get better with use?

5. Read-out

Strong signals: a team with global distribution, very fast shipping, and a hosted moat behind an open SDK. **Key risks:** commoditisation by MCP and model-lab connectors, and the fact that the real moat (reliability at scale) can't be seen in public data. **What would change my view:** per-toolkit reliability metrics and retention cohorts.

Sources: YourStory, "Composio raises $25M led by Lightspeed"; Inc42, "Top 20 Funded AI Startups in India 2026"; GitHub org page github.com/ComposioHQ (stars as of the page fetched 8 Nov 2026); OSS Insight (ComposioHQ/composio); everydev.ai (Composio profile); Merge blog, "Composio pricing" (competitor-written); PyPI JSON API (`composio`, `composio-core`); npm registry (`@composio/core`); Composio docs (framework guides).